



HOUSING PRICE – DATA ANALYSIS AND PREDICTION

ABSTRACT

This paper outlines the data analysis and machine learning methodology to analyze and predict housing price from medium-sized data sample.

By: Zach Nguyen

Contents

Executive summary.....	2
Section 1: Data preparation.....	3
Section 2: Data cleaning, wrangling and Imputation	4
Section 3: Exploratory Data Analysis	8
Section 4: Feature Engineering & Extraction.....	14
Section 5: Data encoding for Machine Learning	21
Section 6: Machine Learning Model Selection and Training	23
Section 7: Conclusion and Risk assessment.....	24

Executive summary

Despite having been around for a long time, Real Estate Price is one of the most unpredictable commodities in the economy. Many buyers and sellers are often unsure of the true value of their properties or those they want to buy. Traditional market research methods such as Similar home or Precedent price comparison introduce significant personal bias and subjectivity. Thus, financial mistakes are myriad in the bidding and negotiation process. This report aims to describe the process of making use of the big data available recently and cutting-edge data science and machine learning methodologies to predict the price of future Real Estate. As another point of reference for property price that is easy to use and interpret, this model can provide Real Estate participants with an advantage in negotiation. To build the model, a sample dataset was taken from Kaggle's "House Prices: Advanced Regression Techniques" and the R Programming language was used as tool of analysis.

Section 1: Data preparation is a minor description of the importing of the sample data and various R packages a practitioner may need when programming and reproducing the methodology. *Section 2: Data cleaning, wrangling and Imputation* describes the data preprocessing phases. Data in real life rarely conforms to ideal expectations of the data scientist. Thus, various techniques to visualize, clean and fill in missing data are conducted. *Section 3: Exploratory Data Analysis* uses visualizations to familiarize the analyst with the structure, state of, and correlation between the dependent variable (Sale Price of the property), in which we hope to predict, and the independent variables, from which we base our predictions. This section also introduces a random forest implementation to grasp important variables that the analyst may need to scrutinize.

Section 4: Feature engineering/Extraction helps convert the survey data into human-sensible data and extract new features (predictors) based on domain understanding. These features reflect real-life concerns of a typical homebuyer looking for a piece of real estate, as surveyed by the analyst. Binary, binning and feature interaction were the techniques used to achieve feature engineering. In *Section 5: Encoding the data to be used in Machine Learning*, the analyst converts the data into machine-sensible data (using dummy variables), reduce any data bias that may prevent the machine from learning efficiently (with normalization) and remove outliers and non-zero-variance data that would introduce unnecessary noise to the model. Finally, *Section 6: Model Selection and Model Training* explains the rationale in selecting a regression model and its modification, coupled with implementation and results obtained by the model.

Overall, you would expect the model to perform with 11.4% mean error in real life. This means that it is overwhelmingly confident that the real value of a \$100,000 property will lie between \$90,000 and \$110,000. The model's regressive nature also introduces interpretability in the form of co-efficient and may help to explain the reason behind the price of certain properties. *Section 7: Conclusion and Risk assessment* will finalize the report and recommend future improvements to the model.

Section 1: Data preparation

This section describes the preliminary data preparation for the rest of the analysis. The data sample was imported from Kaggle.com to build the model. It represents data collected on residential homes in Ames, Iowa. Thus, it is highly advisable for an analyst to find and conduct similar techniques, but on dataset regarding to homes in proximate locations.

The following R packages were used to help with the analysis:

```
16 # Loading libraries
17
18 library(dplyr)           ## Used for Data manipulation, data wrangling
19 library(tidyr)          ## Used for Data manipulation, data wrangling
20 library(VIM)            ## Used to visualization of missing variables
21 library(mice)           ## Used for imputation of missing variables
22 library(psych)          ## Used to get Descriptive statistics
23 library(ggplot2)        ## Used to plot graphs
24 library(gridExtra)      ## Used to arrange visualization
25 library(corrplot)       ## Used for drawing correlation plots
26 library(caret)          ## Used for regression
27 library(randomForest)   ## Used to conduct random forest
28
```

Data was also imported, and non-predictor variables such as ID removed:

```
5 # Importing the data
6 train <- read.csv("train.csv", stringsAsFactors = FALSE)
7 test <- read.csv("test.csv", stringsAsFactors = FALSE)
8
9 test_id <- test$Id
10 test$Id <- NULL
11 train$Id <- NULL
12
13 test$SalePrice <- NA
14 full <- rbind(train,test)
```

Section 2: Data cleaning, wrangling and Imputation

This section depicts the steps to clean and fill in missing data from the dataset. Upon closer look at the dataset, we recognize that some of the data is missing, some are illogical/incorrect, and some are too few and far in between to matter.

We find and clean some illogical data points:

```

50
51 # Find any year value above 2010 and replace with suitable values
52 full %>%
53   select(GarageYrBlt, YrSold, YearBuilt, YearRemodAdd) %>%
54   filter_all(any_vars(.>2010))
55 full[which(full$GarageYrBlt > 2010), "GarageYrBlt"] <- 2007
56
57
58
59 # Find illogical year value (Sold before Built, Remodeled before Built, Garage Built before Built)
60 full %>%
61   select(GarageYrBlt, YrSold, YearBuilt, YearRemodAdd) %>%
62   filter(YrSold < YearBuilt | YearRemodAdd < YearBuilt | GarageYrBlt < YearBuilt)
63 full[which(full$YearBuilt > full$YrSold), "YrSold"] <- 2008
64 full[which(full$YearBuilt > full$YearRemodAdd), "YearRemodAdd"] <- 2002
65
66
67
68 # Find properties with missing garage year built, set them to year built
69 which(!is.na(full$GarageType) & is.na(full$GarageYrBlt))
70 full[c(2127,2577), c("GarageType", "GarageYrBlt", "YearBuilt")]
71 full[c(2127,2577), "GarageYrBlt"] <- full[c(2127,2577), "YearBuilt"]
72

```

We find sparse categories and assign them to the nearest neighbor:

```

68
69 # Find Predictors categories with too few data point and assign to nearest neighbors
70 which(full$MSSubClass == 150)
71 full[2819, "MSSubClass"] <- 160
72
73 which(full$TotRmsAbvGrd == 13 | full$TotRmsAbvGrd == 14 | full$TotRmsAbvGrd == 15)
74 full[c(636,1903,2550), "TotRmsAbvGrd"] <- 12
75
76 which(full$Fireplaces == 4)
77 full[2711, "Fireplaces"] <- 3
78
79 which(full$GarageCars == 5)
80 full[1829, "GarageCars"] <- 4
81


```

In order to tackle missing data, we first visualize the structure of the missing data:

```

42 # Visualize the missingness of the data
43
44 var_with_missing <- colnames(full[colSums(is.na(full)) > 0])
45 var_with_missing <- var_with_missing[-35]
46 aggr(full[,var_with_missing],
47     combined = TRUE, bars = TRUE, sortVars = TRUE, cex.axis = 0.69)
48

```



We also recognize that while some missing data are Missing Completely at Random (MCAR), some missing data such as PoolQC (the variable with the most number data point missing) does not appear random and has a meaning according to the data dictionary provided. This means those are Missing Not at Random (MNAR). To reasonably impute the missing data, we will manually fill in variables with MNAR data and use the MICE package to impute the rest (MCAR).

Missing completely at random

- Using regression can help determine suitable value
- Multiple Imputation by Chained Equation (MICE)

Missing not at random

- Require understanding of missing structure
- Manual imputation required

We fill in the data Missing not at random:

```

84
85 # Group MNAR vars, where NA indicate the subject is not there
86 var_MNAR <- c("Alley", "BsmtQual", "BsmtCond", "BsmtExposure", "BsmtFinType1",
87              "BsmtFinType2", "FireplaceQu", "GarageType", "GarageYrBlt",
88              "GarageFinish", "GarageQual", "GarageCond", "PoolQC", "Fence",
89              "MiscFeature")
90
91 # Create a function to fill in missing data
92 NA_to_NotAvailable <- function(data, predictor) {
93   levels(data[,predictor]) <- c(levels(data[,predictor]), "Not Available")
94   data[,predictor][is.na(data[,predictor])] <- "Not Available"
95   return(data[,predictor])
96 }
97
98 for (i in 1:length(var_MNAR)) {
99   full.MNAR[,var_MNAR[i]] <- NA_to_NotAvailable(full, var_MNAR[i])
100 }
101


```

And visualize the data after the imputation:

```

101
102 # Visualize data after MNAR variables imputation
103 agrgr(full.MNAR[,var_with_missing],
104       combined = TRUE, bars = TRUE, sortVars = TRUE, cex.axis = 0.69)
105

```



It seems that quite a lot of the missing data has been dealt with. Now we use M.I.C.E to impute the remaining missing data points.

```

112
113 # Group the remaining MAR variables
114 var_MAR <- c("MasVnrType", "MSZoning", "Utilities", "Functional",
115             "Exterior1st", "Exterior2nd", "Electrical",
116             "KitchenQual", "SaleType")
117
118 # Convert these into factors for mice imputation
119 full.MNAR <- full.MNAR %>%
120   mutate_at(var_MAR, funs(factor(.)))
121
122 # We use cart method since imputation is fairly linear
123 imputed.data <- mice(full.MNAR[!(names(full.MNAR) %in% "SalePrice")], m=1, method='cart')
124 full.complete <- mice::complete(imputed.data, action = 1, include = FALSE)
125


```

We visualize our imputation to make sure that the values are befitting:

```

125
126 # Visualize Lot Front and Mass Veneer which has the most imputation.
127 # We also use log transformation to linearize graph and aid visualization.
128 xyplot(imputed.data,
129       LotFrontage ~ log10(LotArea), pch = 20, grid = TRUE,
130       xlab = "Area of Lot (Log10 Sq Ft)",
131       ylab = "Length of Street Lot (Ft)",
132       main = "Imputation of Length of Street Lot")
133
134 xyplot(imputed.data,
135       MasVnrArea ~ MasVnrType,
136       xlab = "Type of Mason Veneer",
137       ylab = "Area of Mason Veneer",
138       main = "Imputation of Length of Street Lot")
139

```



The values are quite close, and that's what we want.

```

139
140 # Finalize our data set after imputation
141 train.complete <- full.complete[1:1460,]
142 train.complete$SalePrice <- train$SalePrice
143 test.complete <- full.complete[1460:2919,]
144

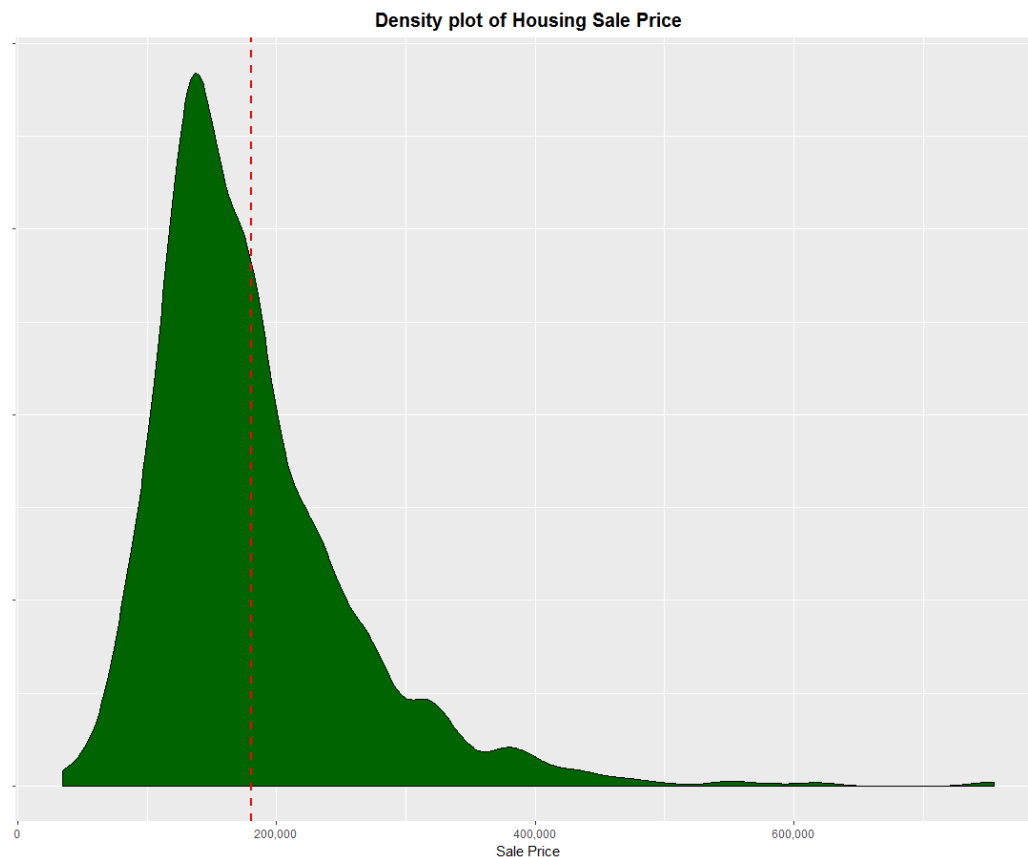
```


Section 3: Exploratory Data Analysis

This section helps the analyst gain intuition on the story the data is telling and aid in feature selection and sanity testing. Having a complete dataset, we can now explore the dataset to gain more understanding and intuition behind the numbers presented. It is, therefore, important to visualize the data structure of the dependent variable, some independent variable that stands out and their correlations.

First, we look at the dependent variable or response variable:

```
154
155 # Visualize the dependent variable SalePrice
156 density_plot_depvar <- ggplot(train.complete, aes(x= SalePrice)) +
157   geom_density(color = "black", fill = "darkgreen") +
158   geom_vline(aes(xintercept = mean(SalePrice)),
159             color = "red", linetype = "dashed", size = 1) +
160   scale_x_continuous(labels = scales::comma) +
161   labs(title = "Density plot of Housing Sale Price", x = "Sale Price", y = "Density") +
162   theme(plot.title = element_text(hjust = 0.5, size = 15, face = "bold"),
163         axis.text.y = element_blank(),
164         axis.title.y = element_blank())
165 density_plot_depvar
166
```



It seems that the average price of a property is around \$175,000. However, many more properties have below average price while the expensive properties “tail off”. This makes sense, as a small percentage of houses is

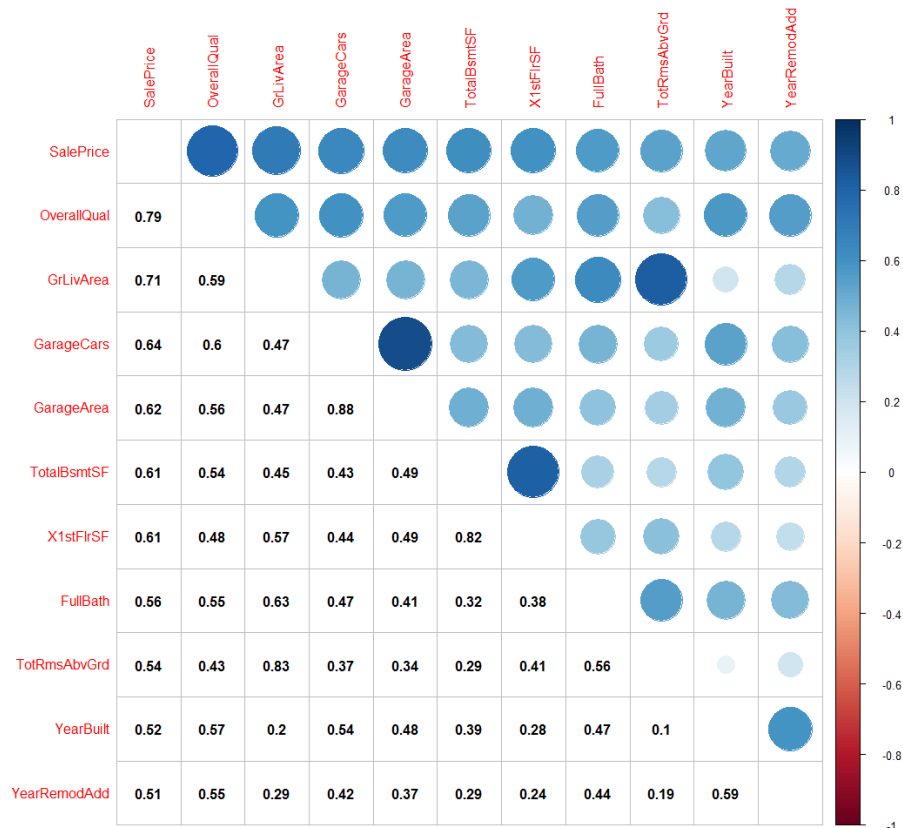
expected to have extraordinary features (mansion, huge land lots, etc. ...) that command much higher price. Therefore, the Sale Price variable may be negatively skewed and will be normalized in later parts.

Next, we'll look at correlations of numeric variables to the Sale Price. This may hint us to some numeric variables that contributes significantly to predicting Sale Price and is, therefore, worth the effort.

```

167 # Visualization of numeric variable correlations to SalePrice
168
169 numeric_var <- which(sapply(train.complete, is.numeric))
170 numeric_var_names <- names(numeric_var)
171
172 # Make correlation chart with all numeric variable
173 cor <- cor(train.complete[,numeric_var], use = "pairwise.complete.obs")
174
175 # Sort chart column in descending order
176 cor_sales_vector <- as.matrix(sort(cor[, "SalePrice"], decreasing = TRUE))
177
178 # Find name index of columns, row with high correlation
179 high_cor_var <- names(which(apply(cor_sales_vector, 1, function(x) abs(x) > 0.5)))
180
181 # Plot correlation chart
182 cor.data <- as.data.frame(cor)[high_cor_var, high_cor_var] # Subset correlation chart
183 cor.data <- cor.data[order(-cor.data$SalePrice),] # Sort row in descending order
184 corplot.mixed(as.matrix(cor.data), lower = "number", lower.col = "black", tl.pos = "lt")
185

```



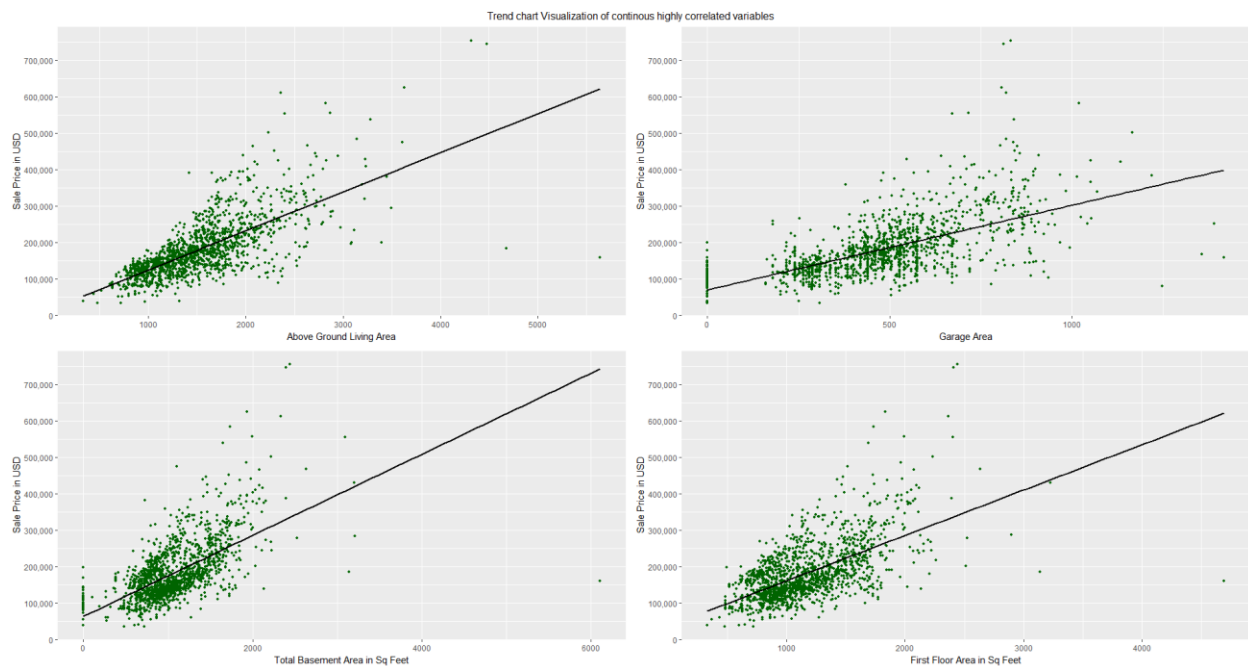
It seems OverallQual score is understandably very highly correlated with SalePrice, along with GrLivArea, 1stFlrSF and TotalBsmtSF. People do prefer bigger houses after all. The YearBuilt and YearRemodAdd variables here also suggest people are looking for newer houses. It is also worth noting that there are highly

correlated variables such as GarageCars and GarageAreas which can cause multicollinearity and bias the model and must be eliminated in latter parts. For now, let's take a closer look at these variables (continuous and discrete) through visualization:

```

185
186 # Visualization of relationships between predictors highly correlated to SalePrice
187
188 Visualize_numeric <- function(data, n_predictor, xname) {
189   ggplot(data, aes(x = data[,n_predictor], y = SalePrice)) +
190     geom_point(na.rm = TRUE, col = "darkgreen", size = 1) +
191     geom_smooth(na.rm = TRUE, method = "lm", se=FALSE, color="black", aes(group=1)) +
192     scale_y_continuous(breaks = seq( 0, max(data$SalePrice), by = 100000), labels = scales::comma) +
193     ylab("Sale Price in USD") +
194     xlab(xname)
195 }
196
197 p1 <- Visualize_numeric(train.complete, "GrLivArea", "Above Ground Living Area")
198 p2 <- Visualize_numeric(train.complete, "GarageArea", "Garage Area")
199 p3 <- Visualize_numeric(train.complete, "TotalBsmtSF", "Total Basement Area in Sq Feet")
200 p4 <- Visualize_numeric(train.complete, "X1stFlrSF", "First Floor Area in Sq Feet")
201 grid.arrange(p1, p2, p3, p4, ncol = 2, top = "Trend chart Visualization of continous highly correlated variables")
202

```



We can definitely see a trend line in the relationship between these variable and Sale Price, the dependent variable. These variables hold promise as important predictors for Sale Price. However, we also notice some outliers.

```

202
203 # We can also spot the outliers
204
205 which(train$GrLivArea > 4500)
206 which(train$GarageArea > 1250)
207 which(train$TotalBsmtSF > 4000)
208 which(train$X1stFlrSF > 4000)
209

```

```

> which(train$GrLivArea > 4500)
[1] 524 1299
> which(train$GarageArea > 1250)
[1] 582 1191 1299
> which(train$TotalBsmtSF > 4000)
[1] 1299
> which(train$X1stFlrSF > 4000)
[1] 1299

```

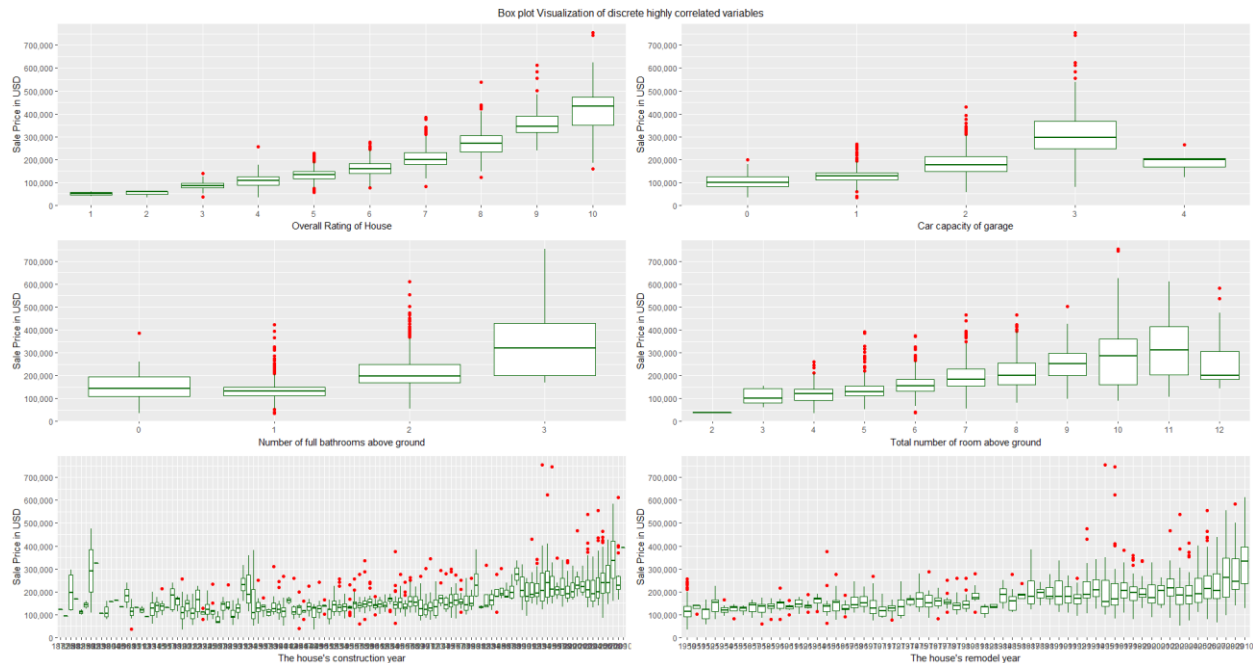
It is necessary to consider removing observation 524 and 1299 at the end. Their low Sale Price despite huge square footage may signify consequence of an imperfect market (e.g.: family, heir transactions, under-the-floor deals ...) or failed negotiation. Though we would generally expect some of these noises to get through our data collection, we can create a better fit model by removing them.

For now, we'll visualize discrete category variables from the correlation plot with Boxplots:

```

209
210 # Visualization of discrete numeric variables
211 Visualize_category <- function(data, c_predictor, xname) {
212   ggplot(data, aes(x= factor(data[,c_predictor]), y = SalePrice)) +
213     geom_boxplot(na.rm = TRUE, col='darkgreen', outlier.color = "red") +
214     scale_y_continuous(breaks= seq( 0, max(data$SalePrice), by = 100000), labels = scales::comma) +
215     ylab("Sale Price in USD") +
216     xlab(xname)
217 }
218
219 p5 <- Visualize_category(train.complete, "OverallQual", "Overall Rating of House")
220 p6 <- Visualize_category(train.complete, "GarageCars", "Car capacity of garage")
221 p7 <- Visualize_category(train.complete, "FullBath", "Number of full bathrooms above ground")
222 p8 <- Visualize_category(train.complete, "TotRmsAbvGrd", "Total number of room above ground")
223 p9 <- Visualize_category(train.complete, "YearBuilt", "The house's construction year")
224 p10 <- Visualize_category(train.complete, "YearRemodAdd", "The house's remodel year")
225 grid.arrange(p5, p6, p7, p8, p9, p10, ncol = 2, top = "Box plot Visualization of discrete highly correlated variables")
226

```



The correlation with SalePrice for these variables can be spotted, especially OverallQual (overall rating). Categories from GarageCars, Full Bath and TotRmsAbvGrd has very discrete boxes (in terms of box height and length) and suggests many areas of differences that could be explored. The Year variables, however, are too spread out (longer boxes) to see a definite trend.

In short, the correlation plot gives us an idea of some classes of important numerical predictors, namely those related to overall quality, garage size and square footage of living spaces. However, we would want to explore important categorical variables as well. Therefore, the plan is to pick up important independent variables through a preliminary random forest implementation.

First, we need to categorize all our independent variables so that R can understand how to feed our data into the machine learning algorithm:

```

228
229 # Specify Factor, ordered factor and numeric values
230 factor_var <- c("MSSubClass", "MSZoning", "Street", "Alley", "LandContour", "LotConfig",
231               "Neighborhood", "Condition1", "Condition2", "BldgType", "RoofStyle", "RoofMatl",
232               "Exterior1st", "Exterior2nd", "MasVnrType", "Foundation", "Heating", "GarageType",
233               "PavedDrive", "MiscFeature", "SaleType", "SaleCondition", "CentralAir", "MSZoning",
234               "HouseStyle", "Fence", "MoSold")
235 ordered_var <- c("LotShape", "Utilities", "LandSlope", "OverallQual",
236                 "OverallCond", "ExterQual", "ExterCond", "BsmtQual",
237                 "BsmtCond", "BsmtExposure", "BsmtFinType1", "BsmtFinType2", "HeatingQC", "Electrical",
238                 "KitchenQual", "Functional", "FireplaceQu", "GarageFinish", "GarageCars",
239                 "GarageQual", "GarageCond", "PoolQC", "GarageYrBlt",
240                 "BedroomAbvGr", "KitchenAbvGr", "TotRmsAbvGrd", "Fireplaces")
241 numeric_var <- c("LotFrontage", "YearBuilt", "YearRemodAdd", "YrSold", "LotArea", "MasVnrArea",
242                 "BsmtFinSF1", "BsmtFinSF2", "BsmtUnfSF", "TotalBsmtSF", "X1stFlrSF", "X2ndFlrSF",
243                 "LowQualFinSF", "GrLivArea", "GarageArea", "WoodDeckSF", "OpenPorchSF",
244                 "EnclosedPorch", "X3SsnPorch", "ScreenPorch", "PoolArea", "MiscVal", "FullBath",
245                 "HalfBath", "BsmtFullBath", "BsmtHalfBath")
246 # Check: length(factor_var) + length(ordered_var) + length(numeric_var) # We have 80 variables
247

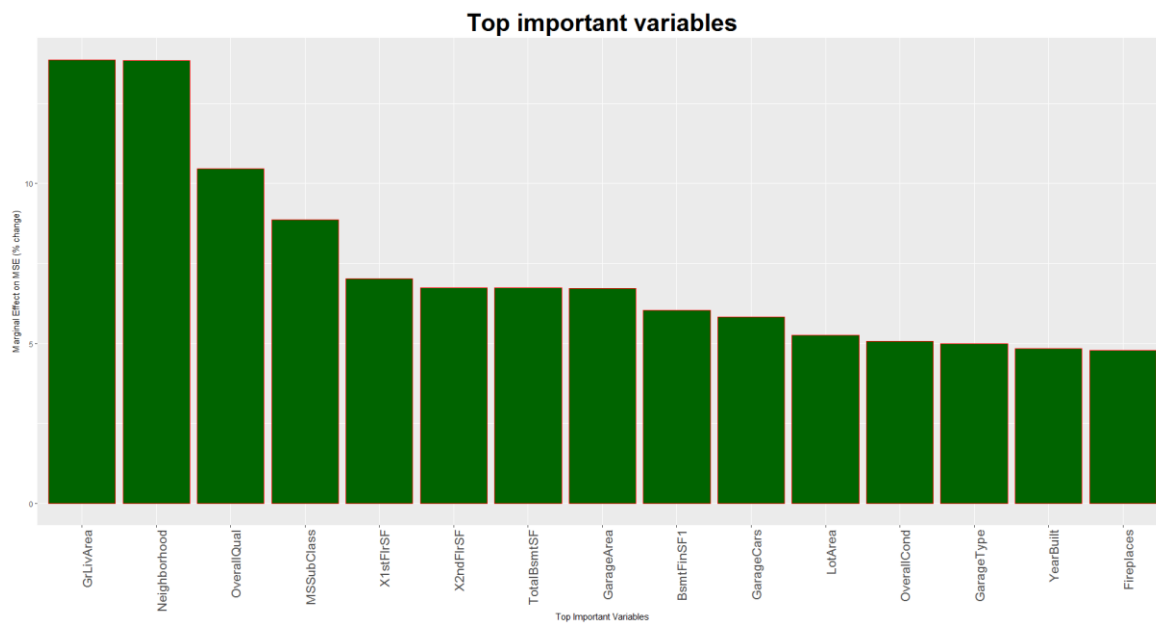
```

Then we train a random forest and use its feature importance methodology to pick out 20% of the most important predictors it considered. The predictors are chosen based on the sensitivity of the model's MSE to changes in said predictors. The 20% choice is a heuristic used only for data exploration, since 20% of predictors usually net 80% of the predictive power of the Machine Learning model.

```

247
248 # Duplicate the dataset to be used in random forest
249 full.rf <- full.complete %>%
250   mutate_at(factor_var, funs(factor(.))) %>%
251   mutate_at(ordered_var, funs(ordered(.)))
252
253 # str(full.rf)
254
255 # Train random forest and pick out 20% of important variables
256 train.rf <- full.rf[1:1460,]
257 train.rf$SalePrice <- train$SalePrice
258
259 set.seed(123)
260 Rf <- randomForest(x = train.rf[, -80], y = train.rf$SalePrice, ntree = 100, importance = TRUE)
261 var.importance <- importance(Rf)
262 var.importance <- data.frame(Variables = row.names(var.importance), MSE = var.importance[, 1])
263 var.importance <- var.importance[order(var.importance$MSE, decreasing = TRUE),]
264 var.importance$Variables <- as.factor(var.importance$Variables)
265
266 # Plot important predictor variables
267 ggplot(var.importance[1:(nrow(var.importance)/5), ], aes(x = reorder(Variables, desc(abs(MSE))), y = MSE)) +
268   geom_bar(stat = "identity", color = "red", fill = "darkgreen") +
269   labs(title = "Top important variables", x = "Top Important Variables", y = "Marginal Effect on MSE (% change)") +
270   theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 1, size = 16),
271         plot.title = element_text(hjust = 0.5, size = 15, face = "bold"))
272

```



Random forest mostly supports our findings from the correlation plot (overall quality, square footage variables, garage and year built). From here, it is clear that most of our efforts should be on these variables. We also discovered important non-numeric categorical variables such as Neighborhood and MSSubClass.


Section 4: Feature Engineering & Extraction

This section describes the feature engineering process. Feature engineering is extremely important in making accurate predictions. Machine Learning models running on insightful data will obtain much better results than those on raw collected data. Due to importance of domain knowledge to feature engineering, I have done market surveys with a sample of 17 of my friends and acquaintances in Toronto to figure out what are the most common questions one would ask when buying a house. Even though bias (demographic, geographic, personal) maybe introduced into the process, it is reasonable to assume that real-estate seekers in North America have more concerns in common than not.

To begin, let us build a function to visualize extractable predictors to see if the common concerns really match the data:

```
276
277 # Build a function to help visualize available variables to extract new sensible ones
278 full.final <- full.complete
279 Visualize_suspect_var <- function(data, s_predictor, xname) {
280   ggplot(data, aes(x = as.factor(data[,s_predictor]), y = data[, "SalePrice"])) +
281     geom_bar(stat = "summary", fun.y = "mean", fill = "green") +
282     geom_bar(stat = "summary", fun.y = "median", fill = "blue", alpha = 0.7) +
283     scale_y_continuous(labels = scales::comma) +
284     geom_hline(yintercept= median(data$SalePrice), linetype="dashed", color = "red", size = 1.5) +
285     geom_hline(yintercept= mean(data$SalePrice), linetype="dashed", color = "orange", size = 1.5) +
286     geom_label(stat = "count", aes(label = ..count.., y = ..count..)) +
287     ylab("Sale Price in USD") +
288     xlab(xname)
289 }
```

Let us also categorize the methods in which we can extract new features from outstanding ones:

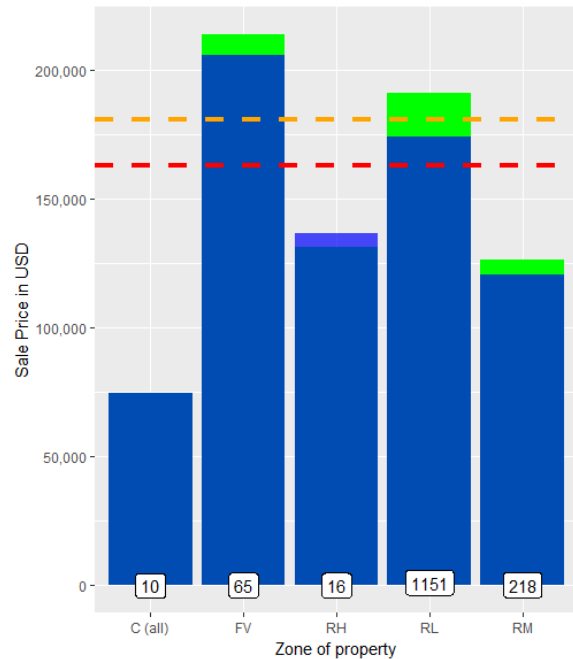


Some new features can be extracted using the survey's recurring questions and the visualization of the data to recognize distinct binary categories. These features describe whether the property is densely populated, new, bought in peak season, has deformed shapes, has detached garage, has paved drive, has been remodelled and bought under a mortgage. These all seem legitimate concerns that would be affect the price of a property. Thus, we devise the categories of the respective variables into reasonable binaries.

```

292
293 # Is the location densely populated?
294 Visualize_suspect_var(train.complete, "MSZoning", "Zone of property")
295 full.final$IsDensePop <- ifelse(full.final$MSZoning %in% c("FV", "RL"), "Yes", "No")
296 full.final$IsDensePop <- as.factor(full.final$IsDensePop)
297


```



```

297
298 # Is the property newer than most?
299 Visualize_suspect_var(train.complete, "MSSubClass", "Class of property")
300 full.final$IsNewer <- ifelse(full.final$MSSubClass %in% c("20", "60", "120", "160"), "Yes", "No")
301 full.final$IsNewer <- as.factor(full.final$IsNewer)
302


```




```

302
303 # Was the property bought during peak season?
304 Visualize_suspect_var(train.complete, "MoSold", "Month Sold of property")
305 full.final$IsPeak <- ifelse(full.final$MoSold %in% c("3","4", "5","6", "7", "8"), "Yes", "No")
306 full.final$IsPeak <- as.factor(full.final$IsPeak)
307


```



```

312
313 # Is lot shape regular, any deformities?
314 Visualize_suspect_var(train.complete, "LotShape", "Lot Shape of property")
315 full.final$IsRegular <- ifelse(full.final$LotShape %in% c("Reg"), "Yes", "No")
316 full.final$IsRegular <- as.factor(full.final$IsRegular)
317


```



```

321
322 # Is Garage Detached?
323 Visualize_suspect_var(train.complete, "GarageType", "Garage Type of property")
324 full.final$IsDetached <- ifelse(full.final$GarageType %in% c("Detchd", "Not Available", "CarPort"),
325                               "Yes", "No")
326 full.final$IsDetached <- as.factor(full.final$IsDetached)
327


```



```

327
328 # Is the Drive Paved?
329 Visualize_suspect_var(train.complete, "PavedDrive", "Drive of property")
330 full.final$IsPaved <- ifelse(full.final$PavedDrive %in% c("Y"), "Yes", "No")
331 full.final$IsPaved <- as.factor(full.final$IsPaved)
332


```



```

332
333 # Has the house been changed/remodeled?
334 Visualize_suspect_var(train.complete, "YearRemodAdd", "Year Remodeled of property")
335 full.final$IsRemod <- ifelse(full.final$YearBuilt != full.final$YearRemodAdd, "Yes", "No")
336 full.final$IsRemod <- as.factor(full.final$IsRemod)
337


```



```

343
344 # Was the property bought under loan/mortgage?
345 Visualize_suspect_var(train.complete, "SaleCondition", "Sales Condition of property")
346 full.final$Invested <- ifelse(full.final$SaleCondition %in% c("Partial"), "Yes", "No")
347 full.final$Invested <- as.factor(full.final$Invested)
348

```



```

342
343 # Does the property have a fireplace?
344 full.final$HasFireplace <- ifelse(full.final$FireplaceQu %in% c("Not Available"), "No", "Yes")
345 full.final$HasFireplace <- as.factor(full.final$HasFireplace)
346
347 # Does the property have a basement?
348 full.final$HasBsmnt <- ifelse(full.final$BsmntQual %in% c("Not Available"), "No", "Yes")
349 full.final$HasBsmnt <- as.factor(full.final$HasBsmnt)
350
351 # Is the property new?
352 full.final$IsNew <- ifelse(full.final$YearBuilt == full.final$YrSold, "Yes", "No")
353 full.final$IsNew <- as.factor(full.final$IsNew)
354
355 # Does the property have a second floor?
356 full.final$Has2ndFlr <- ifelse(full.final$X2ndFlrSF != 0, "Yes", "No")
357 full.final$Has2ndFlr <- as.factor(full.final$Has2ndFlr)
358
359 # Does the property have a Garage?
360 full.final$HasGarage <- ifelse(full.final$GarageArea != 0, "Yes", "No")
361 full.final$HasGarage <- as.factor(full.final$HasGarage)
362
363 # Does the property have a WoodDeck?
364 full.final$HasWD <- ifelse(full.final$WoodDeckSF != 0, "Yes", "No") # Has Wood Deck?
365 full.final$HasWD <- as.factor(full.final$HasWD)
366
367 # Does the property have a Open Porch?
368 full.final$HasOporch <- ifelse(full.final$OpenPorchSF != 0, "Yes", "No") # Has Open Porch?
369 full.final$HasOporch <- as.factor(full.final$HasOporch)
370
371 # Does the property have a Close Porch?
372 full.final$HasCporch <- ifelse(full.final$EnclosedPorch != 0, "Yes", "No") # Has Closed Porch?
373 full.final$HasCporch <- as.factor(full.final$HasCporch)
374
375 # Does the property have a Screen Porch?
376 full.final$HasSporch <- ifelse(full.final$ScreenPorch != 0, "Yes", "No") # Has Screen Porch?
377 full.final$HasSporch <- as.factor(full.final$HasSporch)
378
379 # Does the property feels like its protecting the owners?
380 Visualize_suspect_var(train.complete, "Fence", "Fence of property") # Feels Protected?
381 full.final$Protected <- ifelse(full.final$Fence %in% c("GdPrv", "Not Available"), "Yes", "No")
382 full.final$Protected <- as.factor(full.final$Protected)
383


```

The Neighborhood feature can also be extracted and divided into four distinct categories: Bad, Fair, Good, Excellent. Neighborhood is binned to engineer the NeighborhoodQual feature:

```

387 ggplot(train.complete, aes(x = reorder(Neighborhood, SalePrice, FUN = mean), y = SalePrice)) +
388   geom_bar(stat = "summary", fun.y = "mean", fill = "green") +
389   geom_bar(stat = "summary", fun.y = "median", fill = "blue", alpha = 0.7) +
390   scale_y_continuous(labels = scales::comma) +
391   geom_hline(yintercept= median(train.complete$SalePrice), linetype="dashed", color = "red", size = 1.5) +
392   geom_hline(yintercept= mean(train.complete$SalePrice), linetype="dashed", color = "orange", size = 1.5) +
393   geom_label(stat = "count", aes(label = ..count.., y = ..count..)) +
394   theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 1)) +
395   ylab("Sale Price in USD") +
396   xlab("Neighborhood")
397

```



```

398
399 full.final$NeighborhoodQual <- ifelse(full.final$Neighborhood %in% c("MeadowV"), "Bad",
400                                     ifelse(full.final$Neighborhood %in% c("OldTown", "Edwards", "BrkSide", "Sawyer", "Bluete", "SWISU", "NAmes", "NPKVill"), "Fair",
401                                     ifelse(full.final$Neighborhood %in% c("Mitchel", "SawyerW", "Gilbert", "NWAmes",
402                                     "Blimgin", "CollgCr"), "Good", "Excellent")))
403 full.final$NeighborhoodQual <- as.ordered(full.final$NeighborhoodQual) # Categorize neighborhood by cluster visually
404

```

The following features are also extracted using some of the survey's recurring questions. Contrary to its real value, the property price is often determined by subjective human generalizations. Therefore, we may be able to more accurately predict price if we incorporate some aggregate features such as total bathrooms, total living space and house age.


```

406
407 # Aggregating some features to better reflect human decision making mentality
408 full.final$TotBath <- full.final$FullBath + (full.final$HalfBath*0.5) +
409                     full.final$BsmFullBath + (full.final$BsmHalfBath*0.5) # Total number of bathrooms
410 full.final$TotSF <- full.final$GrLivArea + full.final$TotalBsmSF # Total living space
411 full.final$HouseAge <- full.final$YrSold - full.final$YearBuilt # Total house Age
412

```

Section 5: Data encoding for Machine Learning

This section displays the steps to encode the data for machine learning. Now that we have accomplished a reasonable tailoring of the data and derived the final dataset, we will have to convert the features into their machine-readable forms. These include:



- Dropping highly correlated variables: Correlated independent variables can cause multi-collinearity and overfitting in the model, thereby reducing the interpretability and generalizability power of the model (reducing the precision of a predictor's impact on the independent variable).

```

418
419 # Drop variables with lack of information/ duplicate information
420
421 Visualize_suspect_var(train.complete, "Utilities", "Utilities of property") # Very few observation to matter
422
423 cor.high <- cor
424 cor.high[lower.tri(cor.high, diag = TRUE)] <- NA
425 cor.high <- subset(na.omit(data.frame(expand.grid(dimnames(cor.high)), value = c(cor.high))), value > .7)
426 # Subset to find pairs with high correlation
427 # cor.high
428
429 var.drop <- c("YearRemodAdd", "Utilities", "GarageYrBlt", "GarageArea", "GarageCond", "TotalBsmtSF",
430 "TotalRmsAbvGrd", "BsmtFinSF1")
431 full.final <- full.final[,!(names(full.final) %in% var.drop)] # Drop these variables
432
  
```

- Create levels for ordered factors: We create levels for the algorithm to use as dummy variables (with 0 or 1 values for each observation). This is to get around the fact that the regression can only understand numbers and not string categories. Ordered factors will also give the Machine Learning algorithm more information to work with than ordinary factors.

```

432
433 # Convert the remaining into suitable variable type and add correct levels to ordered
434
435 ordered_var <- ordered_var[!ordered_var %in% c("Utilities", "GarageCond", "GarageYrBlt")]
436 full.final <- full.final %>%
437   mutate_at(factor_var, funs(factor(.))) %>%
438   mutate_at(ordered_var, funs(ordered(.)))
439
440 str(full.final[,ordered_var])
441
442 full.final$SlotShape <- ordered(full.final$SlotShape, levels = c("IR2", "IR2", "IR1", "Reg"))
443 full.final$LandSlope <- ordered(full.final$LandSlope, levels = c("Sev", "Mod", "Gtl"))
444 full.final$ExterQual <- ordered(full.final$ExterQual, levels = c("Po", "Fa", "TA", "Gd", "Ex"))
445 full.final$ExterCond <- ordered(full.final$ExterCond, levels = c("Po", "Fa", "TA", "Gd", "Ex"))
446 full.final$BsmtQual <- ordered(full.final$BsmtQual, levels = c("Not Available", "Fa", "TA", "Gd", "Ex"))
447 full.final$BsmtCond <- ordered(full.final$BsmtCond, levels = c("Not Available", "Po", "Fa", "TA", "Gd"))
448 full.final$BsmtExposure <- ordered(full.final$BsmtExposure, levels = c("Not Available", "No", "Mn", "Av", "Gd"))
449 full.final$BsmtFinType1 <- ordered(full.final$BsmtFinType1, levels = c("Not Available", "Unf", "LwQ", "Rec", "BLQ", "ALQ", "GLQ"))
450 full.final$BsmtFinType2 <- ordered(full.final$BsmtFinType2, levels = c("Not Available", "Unf", "LwQ", "Rec", "BLQ", "ALQ", "GLQ"))
451 full.final$HeatingQC <- ordered(full.final$HeatingQC, levels = c("Po", "Fa", "TA", "Gd", "Ex"))
452 full.final$Electrical <- ordered(full.final$Electrical, levels = c("Mix", "FuseF", "FuseA", "SBrkr"))
453 full.final$KitchenQual <- ordered(full.final$KitchenQual, levels = c("Fa", "TA", "Gd", "Ex"))
454 full.final$Functional <- ordered(full.final$Functional, levels = c("Sal", "Sev", "Maj1", "Maj2", "Mod", "Min2", "Min1", "Typ"))
455 full.final$FireplaceQu <- ordered(full.final$FireplaceQu, levels = c("Not Available", "Po", "Fa", "TA", "Gd", "Ex"))
456 full.final$GarageFinish <- ordered(full.final$GarageFinish, levels = c("Not Available", "Unf", "Rfn", "Fin"))
457 full.final$GarageQual <- ordered(full.final$GarageQual, levels = c("Not Available", "Po", "Fa", "TA", "Gd", "Ex"))
458 full.final$PoolQC <- ordered(full.final$PoolQC, levels = c("Not Available", "Fa", "Gd", "Ex"))
459 full.final$TotRmsAbvGrd <- ordered(full.final$TotRmsAbvGrd, levels = c("2", "3", "4", "5", "6", "7", "8", "9", "10", "11", "12", "13", "14", "15"))
460 full.final$Fireplaces <- ordered(full.final$Fireplaces, levels = c("0", "1", "2", "3", "4"))
461
  
```

- Remove outliers: The outliers identified previously may cause lower the model's accuracy and its ability to generalize outside the test set.

```

462
463 # Removing outliers
464 full.final$SalePrice <- full.final$SalePrice
465 full.final <- full.final[-c(524,1299),]
466
  
```

- Create dummy variables: The dummy variables will greatly increase our number of predictors. Therefore, we will also remove some dummy variable with less than 10 positive observations.

```

466
467 ## Preprocess numeric variables
468
469 nums <- unlist(lapply(full.final, is.numeric))
470 full.numeric <- full.final[,nums]
471 full.numeric <- full.numeric[,!(names(full.numeric) %in% "SalePrice")]
472 full.factor <- full.final[,!nums]
473 cat('There are', length(full.numeric), 'numeric variables, and', length(full.factor), 'factor variables')
474
475 # Create dummy variables
476
477 full.dummy <- as.data.frame(model.matrix(~.-1, full.factor))
478 dim(full.dummy)
479
480 # Taking out a few categorical/dummy variables with very little information
481 var.lowinfo <- which(colSums(full.dummy[1:nrow(full.final[!is.na(full.final$SalePrice),]),])<10)
482 colnames(full.dummy[var.lowinfo])
483
484 full.dummy <- full.dummy[, -var.lowinfo] #removing variables
485 dim(full.dummy)
486


```

- Normalization & Log transformation: We observe skewness in some independent variable and the response variable. Therefore, we will log transform these variables. We will also normalize the numeric predictors to center and scale them for the Machine Learning algorithm.

```

474
475 # Log transformation of numeric variables
476 describe(full.numeric)
477
478 for(i in 1:ncol(full.numeric)){
479   if (abs(skew(full.numeric[,i]))>0.8){
480     full.numeric[,i] <- log(full.numeric[,i] +1)
481   }
482 }
483
484 # Normalizing the numeric variables
485 num.process <- preProcess(full.numeric, method=c("center", "scale"))
486 print(num.process)
487 full.numeric.scaled <- predict(num.process, full.numeric)
488
504
505 # Dealing with skewness in response variable
506 skew(full.final$SalePrice)
507 full.encode$SalePrice <- log(full.final$SalePrice)
508 qqnorm(full.encode$SalePrice)
509 qqline(full.encode$SalePrice) # Log transform and visualize SalePrice
510

```



Finally, we combine the two subsets to an encoded dataset. We will use this to derive the encoded training set and the encoded test set.

```

500
501 # Combining the predictor variables
502 full.encode <- cbind(full.numeric.scaled, full.dummy)
503
510
511 # Divide encoded data into train and test set
512 train.encode <- full.encode[1:1458,]
513 test.encode <- full.encode[1459:2917,]
514

```

Section 6: Machine Learning Model Selection and Training

This section attempts to train a suitable machine learning algorithm to predict the sale price of future properties. The problem here is a prediction of continuous value for the response variable. Thus, we can begin with conducting a simple regression. However, with only a few predictors showing strong predictive power out of over 200 features, we will want to introduce a regularization parameter to penalize the number and/or effect of features used by our regression. Thus, we will use an elastic net via cross-validation to best tune our regularization parameters (alpha as degree of mix between LASSO and Ridge regression and lambda as the shrinkage parameter of the penalty).

```

518
519 set.seed(123)
520 mod.control <- trainControl(method="cv", number=5)
521 grid <- expand.grid(alpha = 1, lambda = seq(0.001,0.1,by = 0.001))
522
523 regressor <- train(x=train.encode[,!(names(train.encode) %in% "SalePrice")], y=train.encode$SalePrice,
524 | | | | | method='glmnet', trControl= mod.control, tuneGrid= grid)
525 regressor$bestTune
526

```

```

alpha lambda
2      1 0.002

```

Our hyperparameter tuning shows best results for LASSO regression with a shrinkage parameter of 0.002. This will also result in a RMSE score of ~ **11.4%**:

```

> min(regressor$results$RMSE)
[1] 0.1138914

```

Let's also check if Lasso Regression did its job by cutting down on many of the variables.

```

529
530 # Check the number of variables being selected for the best model.
531 contrib <- varImp(regressor,scale=F)$importance
532 var.selected <- length(which(contrib$Overall!=0))
533 var.nselected <- length(which(contrib$Overall==0))
534
535 cat('L-Regression uses', var.selected, 'variables in its model, and did not select', var.nselected, 'variables.')
536

```

```
L-Regression uses 110 variables in its model, and did not select 128 variables.
```

It has removed 128 variables and used 110. The regularized regression seems to be using only independent variables with statistical significant. It therefore looks to be a viable model.

Section 7: Conclusion and Risk assessment

Throughout this report, we have prepared, cleaned and explored the dataset collected. Then, we engineered new features with domain understanding to help our predictive model perform better and chose a suitable machine learning model for the prediction task. Our model is likely to score 11.4% precision with any test set and has been rigorously devised to generalize well in reality.

However, improvements can certainly be investigated. The following are risks and improvements that users of this report can make to improve the model's performance:

- First, the user should care to use a training dataset that accurately reflects the geographic, social and demographic trends relevant to real estate transactions in the setting of the consulting property.
- Second, the user should explore more algorithm options. Although simple linear regression models are usually best when predicting a generally linear problem such as this, results could be enhanced with incorporation of non-linear terms. If accuracy over speed is the user's primary concern, he/she should explore using neural networks, gradient boosting trees or other combination and ensemble of models to derive predictions.
- Lastly, the user should only consider the results of the algorithm as a referred generalized price point. Some rare properties and real-estates can have characteristics that make them extra-ordinary and will most likely not be accounted for in the model.